

Learning the latent spaces dedicated to the segmentation of medical imaging

Application to cardiac imaging

Pr. Olivier Bernard

Lab. CREATIS – Univ. of Lyon, France

olivier.bernard@creatis.insa-lyon.fr

CREATIS



What is the interest of generative models ?


► How to generate synthetic faces ?



By modeling the corresponding distribution $p_\theta(\cdot)$!

What are the interest of generative models ?

► How to model complex distributions ?




Face distribution

What are the interest of generative models ?

► What for ?

One obsession is to master the latent space !!!

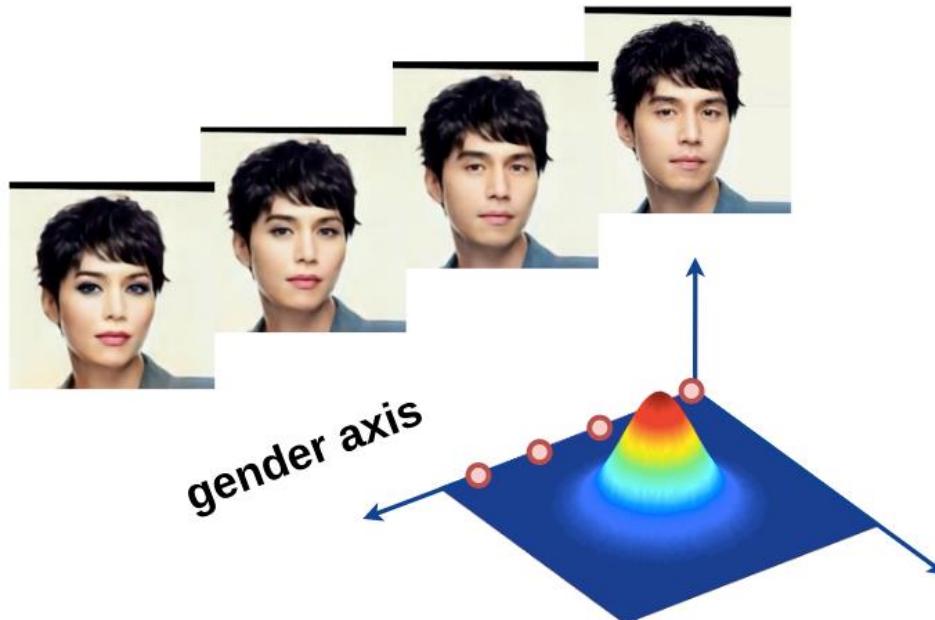


Latent space $z_i \in \mathbb{R}^K$

What are the interest of generative models ?

► What for ?

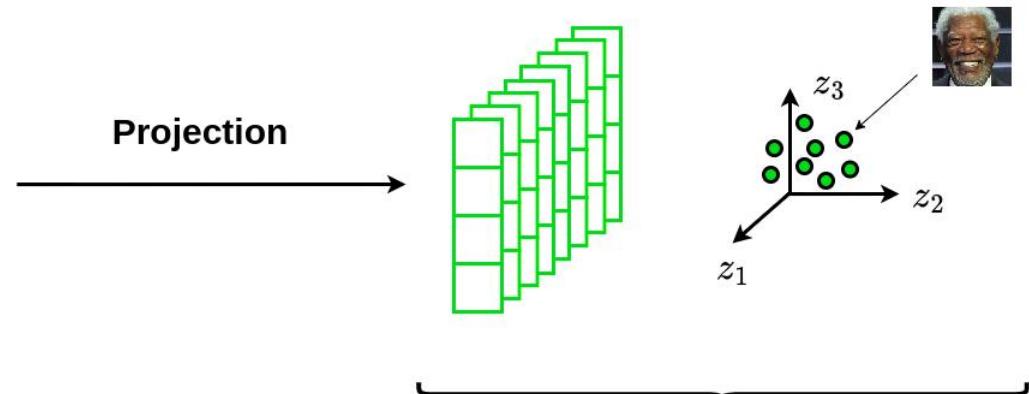
One obsession is to master the latent space !!!



Auto-encoders

How to learn a distribution ?

- ▶ Projection into a simpler, lower-dimensional representation space




Latent space $z_i \in \mathbb{R}^K$

Input space $x_i \in \mathbb{R}^{N \times M}$

How to learn a complex distribution ?

- ▶ How to have a relevant representation space ?



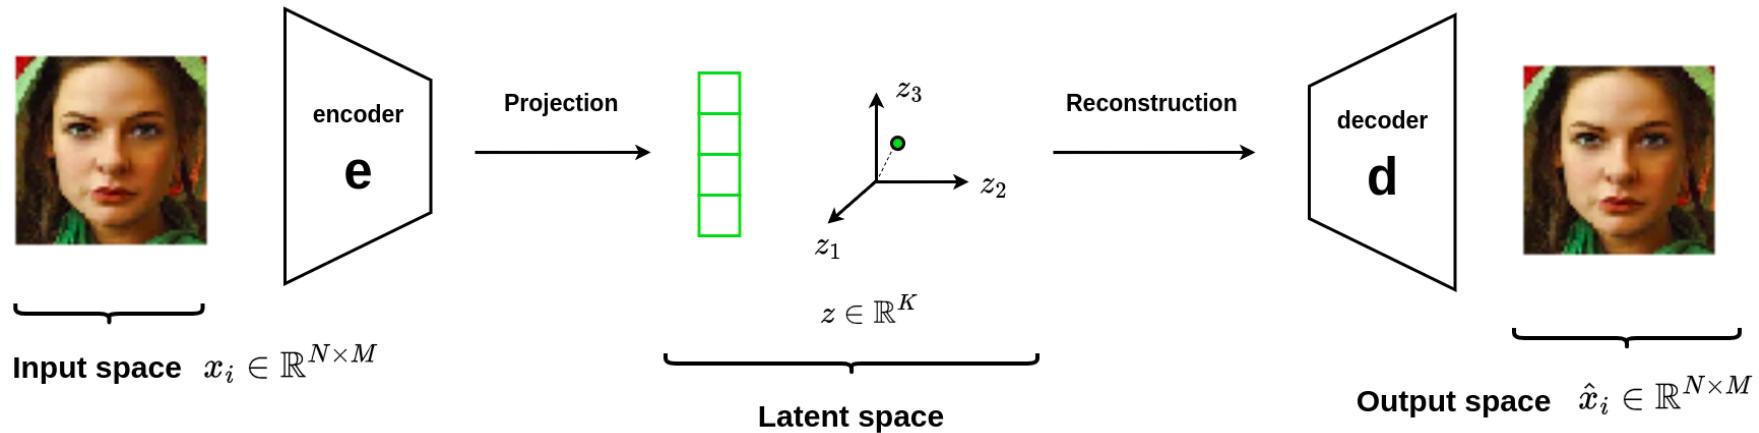
Reconstruction →



Output space $\hat{x}_i \in \mathbb{R}^{N \times M}$

Auto-encoder framework

► Standard architecture





► Deep learning loss function

$$\text{loss} = \|x - \hat{x}\|^2$$


Interest of auto-encoders

► Generative model

0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9




0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9



Limitations


- ▶ Needs to better control the structure of the latent space




Interest of auto-encoders

- Generative model with better properties thanks to *variational framework*


0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9



2nd component



1st component




0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9

Latent space $z_i \in \mathbb{R}^2$


Interest of auto-encoders

► Generative model with variational framework



Linear interpolation into the latent space

$$t \cdot z_0 + (1 - t) \cdot z_7, \quad 0 \leq t \leq 1$$



Variational autoencoders

All the mathematical details are given there !

<https://creatis-myriad.github.io/tutorials/2022-09-12-tutorial-vae.html>

Key concepts

- ▶ Enforcing a structured latent space



 - Through a probabilistic framework

 - By imposing continuity

 - By imposing completeness

Probabilistic framework

► Mathematical formulation



Approximation of $p(z|x)$ through a variational inference technique


Probabilistic framework

► Hypotheses

→ $q(z|x)$ is modeled by an axis-aligned Gaussian distribution

→ $q(z|x) = \mathcal{N}(\mu_x, \sigma_x) = \mathcal{N}(g(x), \text{diag}(h(x)))$

$$(g^*, h^*) = \arg \min_{(g,h)} D_{KL}(q(z|x) \parallel p(z|x))$$



$D_{KL}(\cdot \parallel \cdot)$ Kullback-Liebler divergence function

Probabilistic framework

► Optimization process


→ Maximization of the Evidence Lower Bound (ELBO)

$$\mathcal{L} = \mathbb{E}_{z \sim q_x} [\log(p(x|z))] - D_{KL}(q(z|x) \parallel p(z))$$

→ By exploiting gaussian assumption

$$p(x|z) = \mathcal{N}(f(z), cI)$$

$$\mathcal{L} \propto \mathbb{E}_{z \sim q_x} [-\alpha \|x - f(z)\|^2] - D_{KL}(q(z|x) \parallel p(z))$$




Probabilistic framework

► Optimization process

$$(f^*, g^*, h^*) = \arg \min_{(f,g,h)} \left(\mathbb{E}_{z \sim q_x} [\alpha \|x - f(z)\|^2] + D_{KL}(q(z|x) \parallel p(z)) \right)$$

► Deep learning loss function

$$\text{loss} = \alpha \|x - f(z)\|^2 + D_{KL}(\mathcal{N}(g(x), \text{diag}(h(x))), \mathcal{N}(0, I))$$




→ $g(\cdot)$ and $h(\cdot)$ are modeled through an encoder

→ $f(\cdot)$ is modeled through a decoder

Probabilistic framework

► Loss interpretation


$$\text{loss} = D_{KL} (\mathcal{N}(g(x), \text{diag}(h(x))), \mathcal{N}(0, I)) + \alpha \|x - f(z)\|^2$$



Probabilistic framework

► Loss interpretation

$$\text{loss} = D_{KL} (\mathcal{N}(g(x), \text{diag}(h(x))), \mathcal{N}(0, I)) + \alpha \|x - f(z)\|^2$$





- $\mathcal{N}(g(x), h(x))$ imposes local **continuity**
- $\mathcal{N}(\cdot, \mathcal{N}(0, I))$ imposes global **completeness**

Practical applications

The obsession is to master the latent space !!!


Needs for accurate and robust segmentation of cardiac structures

- ▶ Quantification of clinical indices from echocardiographic images



Needs for accurate and robust segmentation of cardiac structures

► Anatomical clinical indices




How to guarantee the anatomical coherence ?

► Constrained Variational Auto Encoder


- Approximation of a latent space with local linear properties


Use of a 1-neuron net to reinforce the linearity of the latent space




► Constrained Variational Auto Encoder

- Approximation of a latent space with local linear properties
 - ➔ Linear interpolation in the latent space makes sense







Efficient encoding of anatomical shapes in a latent space



► Definition of 12 anatomical metrics

- (3 criteria) hole(s) in the LV, RV or LA
- (2 criteria) hole(s) between LV and MYO or between LV and LA
- (3 criteria) presence of more than one LV, MYO or LA
- (2 criteria) size of the area by which the LV touches the background or the MYO touches the LA
- (1 criterion) ratio between the minimal and maximal thickness of the MYO
- (1 criterion) ratio between the width of the LV and the average thickness of the MYO






Correction of segmentation to guarantee the plausibility of anatomical shapes




Almost same accuracy as the original methods but with correct anatomical shapes




How to guarantee temporal consistency ?

- ▶ Quantification of clinical indices from echocardiographic images

What we have with a 2D U-Net




What we want




► AR-VAE: attribute-based regularization of VAE latent space
[Pati, Neural Comp. Appl., 2021]

- Generation of structured latent space
 - Specific continuous-valued attributes forced to be encoded along specific dimensions
 - $\text{Loss} = \text{VAE loss} + \text{Attribute Regularisation Loss}$



► AR-VAE: attribute-based regularization of VAE latent space
[Pati, Neural Comp. Appl., 2021]




- Sampling of the structured latent space



► AR-VAE: attribute-based regularization of VAE latent space


[Pati, Neural Comp. Appl., 2021]


- Sampling of the structured latent space
 - Specific attribute (from left to right): area, length, thickness, slant, width, height
 - Each column corresponds to traversal along a regularized dimension




► Application to the description of the cardiac shapes

- Generation of structured latent space according to the following attributes
 - ➔ Left ventricle (LV) cavity: area, length, basal width, orientation
 - ➔ Myocardial area
 - ➔ Epicardial center






► Proposed temporal pipeline




Cardiac segmentation with temporal consistency

[Painchaud, IEEE TMI, 2022]




- ▶ Some post-processing examples

Original U-Net




Post-processed U-Net




- ▶ Some post-processing examples


Original U-Net




Post-processed U-Net




Uncertainty estimation for cardiac image segmentation



Inference procedure



► Uncertainty results



To conclude

To conclude


- ▶ VAEs can be used effectively in medical imaging
 - Guarantee anatomical coherence ✓
 - Guarantee temporal consistency ✓
 - Estimation uncertainty for image segmentation ✓
 - Generative interest limited to simple distribution
- ▶ Useful tool for characterizing populations
 - Need to properly structure the learned latent space
 - Need to work on relatively large cohorts

Appendix

Probabilistic framework

► Continuity


$$\mathcal{N}(g(x), h(x))$$



Probabilistic framework


► Completeness

$$\mathcal{N}(\cdot, \mathcal{N}(0, I))$$




Deep learning implementation

► Reparameterization trick




Interest of auto-encoders


► Data representation



► Temporal inconsistency detection from the latent space





Choppy contraction/dilation of
the LV cavity



Abrupt vertical shifts of the
cardiac shape


► Rejection sampling

- Targeted distribution $P(z)$
 - Parzen window technique
- Proposed distribution $Q(z)$
- Constrain $kQ(z) > P(z)$
 - Automatic choice of k





► Rejection sampling

- $\mathbf{z} \sim Q(\mathbf{z})$
- $\mathbf{u} \sim \text{Unif}(0, kQ(\mathbf{z}))$
- Computation of $P(\mathbf{z})$
 - If $u \leq P(\mathbf{z})$ then keep \mathbf{z}
 - If $u > P(\mathbf{z})$ then reject \mathbf{z}




Needs for accurate and robust segmentation of cardiac structures

► Functional clinical indices



- Volume dynamic of the cavities over the cardiac cycle
- Global longitudinal strain of the heart muscle




► Quantitative evaluation

- CAMUS dataset

- 500 patients x 2 probe orientation x 2 key frames
- 2000 images with reference contours
- Metrics: Dice / Hausdorff dist.

Example of a segmentation result



Methods	Original	VAE		Nearest Neighbors	
	-	Robust	w/o RS	w/ RS	
U-Net [5], [8]	.921 / 6.0	.923 / 5.7	.923 / 5.7	.922 / 5.7	.922 / 5.7
LUNet [14]	.922 / 5.9	.921 / 5.9	.922 / 5.9	.921 / 5.9	.921 / 6.0
ENet [31]	.923 / 5.8	.921 / 5.9	.921 / 5.9	.920 / 5.9	.920 / 5.9
SHG [32]	.915 / 6.2	.915 / 6.2	.916 / 6.2	.915 / 6.2	.915 / 6.2
SRF [33]	.879 / 13.1	.877 / 13.2	.878 / 13.2	.879 / 13.0	.879 / 13.0
BEASM-auto [34], [35]	.868 / 10.5	.868 / 10.5	.867 / 10.5	.868 / 10.5	.868 / 10.5
BEASM-semi [5], [34]	.899 / 7.8	.899 / 7.8	.899 / 7.8	.899 / 7.8	.899 / 7.8

► Quantitative evaluation

- Corr: Correlation between the sum of the uncertainty values (foreground) and Dice score
- MI: Mutual Information between the uncertainty map and the error map

Method	Training data		CAMUS		Shenzen	
	Corr. ↑	MI ↑	Corr. ↑	MI ↑	Corr. ↑	MI ↑
Entropy	0.66	0.02	0.34	0.02	0.89	0.02
ConfidNet [1]	0.34	0.04	0.36	0.04	0.69	0.01
<i>CRISP</i>	0.71	0.20	0.41	0.06	0.83	0.11
McDropout [3]	0.67	0.03	0.26	0.02	0.82	0.03
<i>CRISP-MC</i>	0.78	0.26	0.29	0.06	0.82	0.08
LCE [2]	0.58	0.44	0.35	0.37	0.87	0.06
<i>CRISP-LCE</i>	0.59	0.08	0.34	0.13	0.85	0.11