

Generative AI for medical imaging

Recent advances and potential benefits




Olivier Bernard



olivier.bernard@insa-lyon.fr

Generative AI for imaging

Generative AI for medical imaging



Generative AI for medical imaging

► Key challenges

Generative
ability

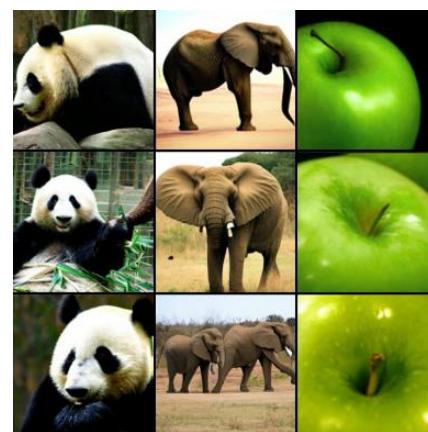
Conditioning

Multimodality

Real images



Synthetic images



An Asian girl in ancient
rides a giant panda

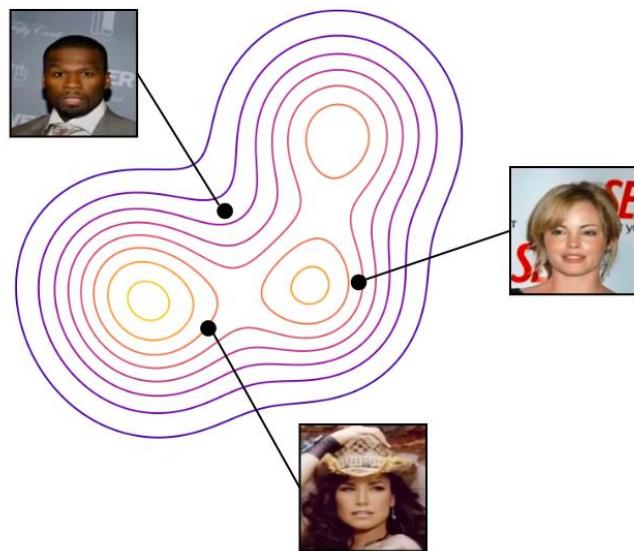


Generative AI for medical imaging

Generative ability

Real distribution

$$p(x)$$



Learning




Sampling



Synthetic images

Generative AI for medical imaging



Generative AI for medical imaging

Multimodality

An Asian girl in ancient coarse linen clothes rides a giant panda and carries a wooden cage. A chubby little girl with two buns walks on the snow. High-precision clothing texture, real tactile skin, foggy white tone, low saturation, retro film texture, tranquil atmosphere, minimalism, long-range view, telephoto lens




Generative AI for medical imaging

Multimodality

Text-to-Video

* Hover over the video to see the input prompt.



https://lumiere-video.github.io/#section_image_to_video

Generative AI for medical imaging

Multimodality

Image-to-Video

* Hover over the video to see the input image and prompt.





https://lumiere-video.github.io/#section_image_to_video

Generative AI for medical imaging

► Family of networks


VAE

sampling





GAN

sampling



Diffusion
models


sampling



Variational auto-encoders

How to learn a complex distribution ?

- ▶ Projection of the data into a **lower dimensional space** called latent space
- ▶ Interest: generating a more **compact and interpretable representation**




Latent space $z_i \in \mathbb{R}^K$

Input space $x_i \in \mathbb{R}^{N \times M}$


How to learn a complex distribution ?

- ▶ How to learn a relevant latent representation ?



Auto-encoder framework

► Standard encoder / decoder architecture



► Deep learning loss function

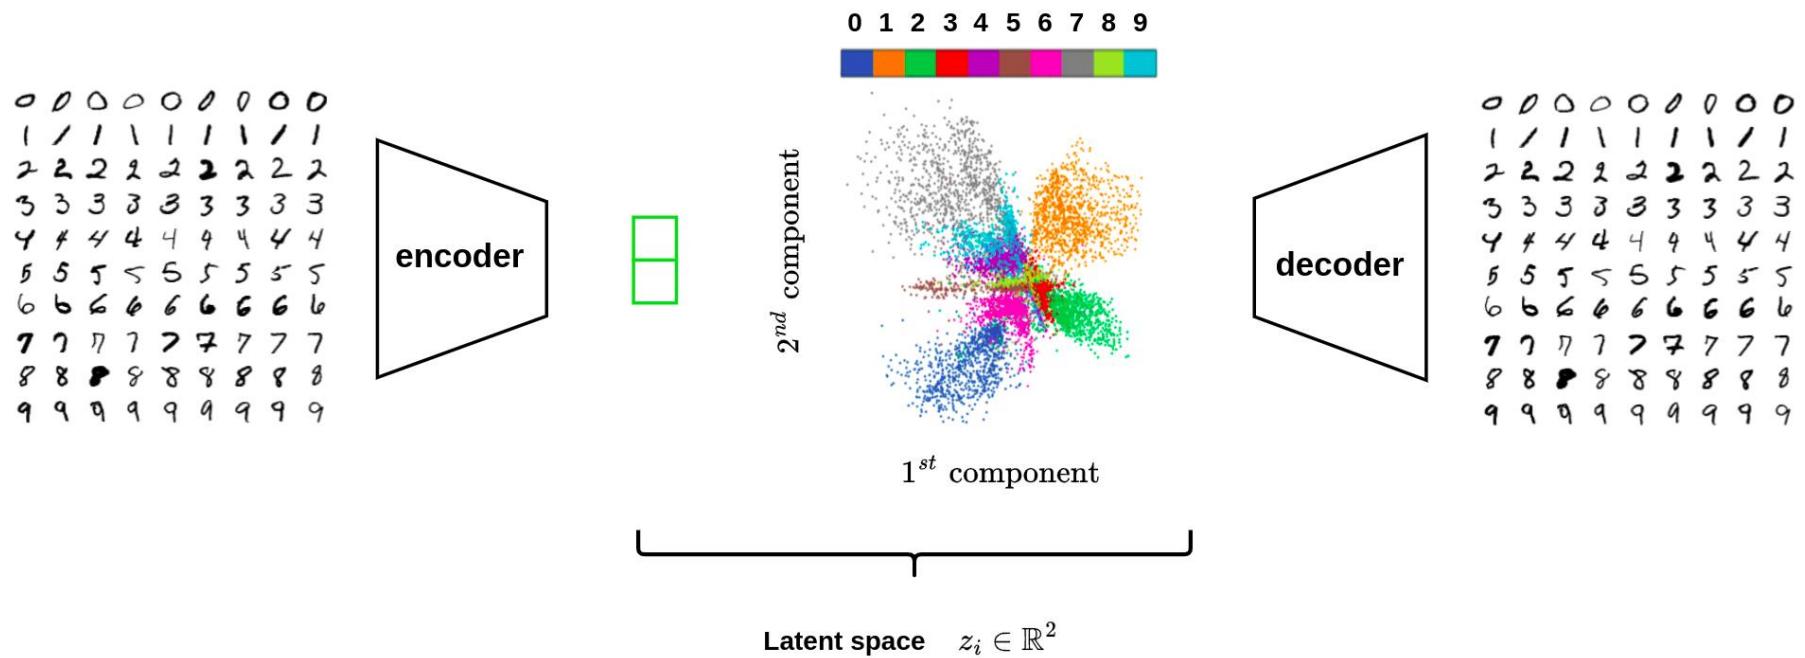
$$\text{loss} = \|x - \hat{x}\|^2$$

$$\text{loss} = \|x - d(e(x))\|^2$$

Auto-encoder weaknesses


► Illustration from MNIST dataset

- (train,test) = (60 000, 10 000)
- Input image size: 32x32 / latent space $K=16$ (compression factor around 64)





Auto-encoder weaknesses

- ▶ Needs to better control the structure of the latent space



Auto-encoder weaknesses


- ▶ Sampling random latent vector



Variation Auto Encoder framework

► Starting point


Intractable distribution



Data distribution

$$x \sim p(X) \in \mathbb{R}^{N \times M}$$

Controlled distribution




Latent distribution

$$z \sim p(Z) \in \mathbb{R}^K$$

Variation Auto Encoder framework


- ▶ Creating a mapping between the two distributions
 - ➔ Through Bayesian statistics



Variational Auto Encoder framework

► Strong assumptions

- Latent distribution $p(z)$ is assumed to be a normal distribution
- The likelihood distribution is $p(x|z)$ assumed to be a Gaussian distribution whose parameters need to be learned
- The posterior distribution $p(z|x)$ is intractable and needs to be approximated



- ▶ Approximation of the posterior through variational inference
 - Statistical approximation technique for complex distributions, here $p(z|x)$
 - Definition of a parameterized family of distributions
 - ▶ e.g., family of Gaussian distributions with parameters μ_x, σ_x modeled by functions to be determined
 - Find the best approximation of the target distribution in this family
 - The best element of the family minimizes an approximation error measure between two distributions
 - ▶ Kullback-Leibler divergence function is often used

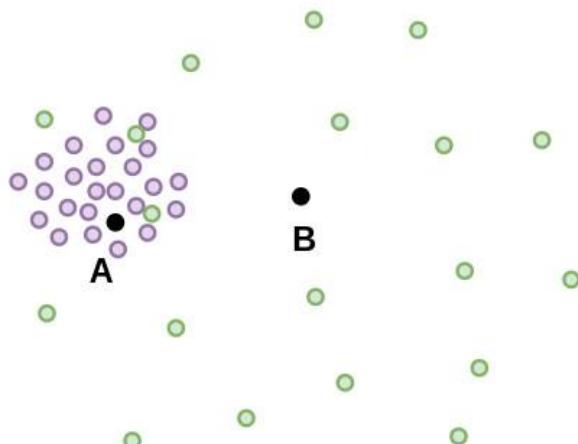
► Kullback-Leibler divergence function

→ Distance measure between two distributions via relative entropy

$$D_{KL}(p \parallel q) = \int p(x) \cdot \log \left(\frac{p(x)}{q(x)} \right) dx$$

→ D_{KL} is a measure that is always positive $D_{KL}(p||q) \geq 0$


→ D_{KL} is a nonsymmetric measure $D_{KL}(p||q) \neq D_{KL}(q||p)$



- For the purple distribution, the distance AB is large
- For the green distribution, the distance AB is moderate
- The notion of distance differs depending on the distributions

Variational Auto Encoder framework

- ▶ Enforce a structured latent space with reduced dimensionalities
 - Through Bayesian statistics



Variation Auto Encoder framework

- ▶ The prior is modeled through a Gaussian distribution

$$p(z) = \mathcal{N}(0, I)$$

- ▶ The likelihood is modeled through a Gaussian distribution

$$p(x | z) = \mathcal{N}(\mu_z, \sigma_z) = \mathcal{N}(f(z), cI)$$


- ▶ The posterior is approximated by an axis-aligned Gaussian distribution

$$q(z | x) = \mathcal{N}(\mu_x, \sigma_x) = \mathcal{N}(g(x), diag(h(x)))$$

Variation Auto Encoder framework

► Optimization process

- Sample a new data from the original data distribution
- Pick a sample x that maximize $p(x)$, or $\log(p(x))$



$\log(p(x))$

$$\log(p(x)) = \log \left(\int p(x, z) dz \right)$$



Marginal distribution

$$\log(p(x)) = \log \left(\int \frac{q(z|x)}{q(z|x)} p(x, z) dz \right)$$



Expectation
reformulation

$$\log(p(x)) = \log \left(\mathbb{E}_{q(z|x)} \left[\frac{p(x, z)}{q(z|x)} \right] \right)$$




Jensen's inequality

$$\log(p(x)) \geq \mathbb{E}_{q(z|x)} \left[\log \left(\frac{p(x, z)}{q(z|x)} \right) \right]$$

Variational Auto Encoder framework

► Evidence lower bound (ELBO)

$$\log(p(x)) \geq \mathbb{E}_{q(z|x)} \left[\log \left(\frac{p(x, z)}{q(z|x)} \right) \right] \quad \text{ELBO}$$



→ Maximization of the ELBO

$$\mathcal{L}(x) = \mathbb{E}_{q(z|x)} \left[\log \left(\frac{p(x, z)}{q(z|x)} \right) \right]$$

$$\mathcal{L}(x) = \mathbb{E}_{q(z|x)} \left[\log \left(\frac{p(x|z)p(z)}{q(z|x)} \right) \right]$$

$$\mathcal{L}(x) = \mathbb{E}_{q(z|x)} \left[\log(p(x|z)) + \log \left(\frac{p(z)}{q(z|x)} \right) \right]$$

$$\mathcal{L}(x) = \mathbb{E}_{q(z|x)} [\log(p(x|z))] - D_{KL}(q(z|x) \| p(z))$$


Bayes' formula
 $p(x, z) = p(x|z)p(z)$

Kullback-Liebler divergence D_{KL}

Variation Auto Encoder framework

► ELBO maximization

$$\mathcal{L} = \mathbb{E}_{z \sim q_x} [\log(p(x|z))] - D_{KL}(q(z|x) \parallel p(z))$$



→ Exploitation of the Gaussian assumption of the likelihood


$$p(x|z) = \mathcal{N}(f(z), cI)$$

$$\mathcal{L} \propto \mathbb{E}_{z \sim q_x} [-\alpha \|x - f(z)\|^2] - D_{KL}(q(z|x) \parallel p(z))$$

Variational Auto Encoder framework

► Optimization process

$$(f^*, g^*, h^*) = \arg \min_{(f,g,h)} (\mathbb{E}_{z \sim q_x} [\alpha \|x - f(z)\|^2] + D_{KL}(q(z|x) \parallel p(z)))$$



► Deep learning loss function

$$\text{loss} = \alpha \|x - f(z)\|^2 + D_{KL}(\mathcal{N}(g(x), \text{diag}(h(x))), \mathcal{N}(0, I))$$

→ $g(\cdot)$ and $h(\cdot)$ are modeled through an encoder


→ $f(\cdot)$ is modeled through a decoder

Variation Auto Encoder framework

► Interpretation of the loss function


$$\text{loss} = \alpha \|x - f(z)\|^2 + D_{KL} (\mathcal{N}(g(x), \text{diag}(h(x))), \mathcal{N}(0, I))$$

Data attachment term *Completeness* constraint *Continuity* constraint



Variational Auto Encoder framework

► Reparameterization trick





Variational framework


► Illustration from MNIST dataset

- (train,valid,test) = (50 000,10 000,10 000)
- Input image size: 28x28 / latent space $K=16$ (compression factor around 50)

0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9




0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9




Variational framework

► Generative model with variational framework



Linear interpolation into the latent space

$$t \cdot z_0 + (1 - t) \cdot z_7, \quad 0 \leq t \leq 1$$



Reinforcement of the generative process

Structuration of the latent space: AR-VAE

► Structuration of latent space based on image attributes


- What is an attribute ?

- ➔ Measurement performed in image space to characterize a target object

- ➔ E.g.: handwritten digits (MNIST database)


- ▶ Attributes: line thickness, inclination, length, area, ...

- ➔ Pre-training image attribute measurements used as input data



Structuration of the latent space: AR-VAE

- ▶ Structuration of latent space based on image attributes
 - Each attribute are coded according to a specific latent dimension



Structuration of the latent space: AR-VAE

► Attribute regularization term

- During the learning phase

➔ Computation for each attribute a of a distance matrix $D_a \in \mathbb{R}^{m \times m}$ from the m images $\{x_i\}_{1 \leq i \leq m}$ present in the current batch

$$D_a(i, j) = a(x_i) - a(x_j) \quad \text{with} \quad i, j \in [0, m)$$

➔ Computation for each attribute r of a distance matrix $D_r \in \mathbb{R}^{m \times m}$ from the m latent vector $\{z_i\}_{1 \leq i \leq m}$ corresponding to the images in the current batch

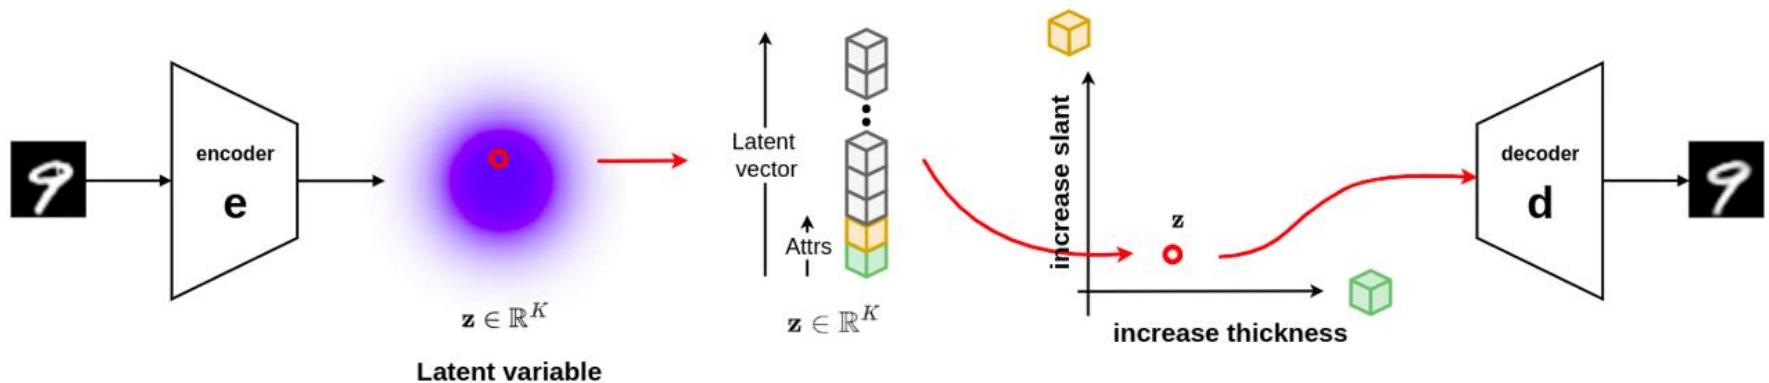
$$D_r(i, j) = z_i^r - z_j^r \quad \text{with} \quad i, j \in [0, m)$$

➔ Introduction of the following loss term

$$\mathcal{L}_{r,a} = MAE(\tanh(D_r) - sign(D_a))$$


Structuration of the latent space: AR-VAE

- Generate a latent space structured according to attributes



Structuration of the latent space: AR-VAE


- ▶ Generate a latent space structured according to attributes
 - Sampling of the structured latent space



Structuration of the latent space: AR-VAE

- ▶ Generate a latent space structured according to attributes

- Sampling of the structured latent space
 - Specific attributes: surface, length, thickness, inclination, width, height
 - Each column corresponds to a traverse along a regularized dimension




Medical applications


Shape representation

► Application example: representation of cardiac shapes

- Generation of a latent space structured according to the following attributes
 - ➔ Left ventricular (LV) cavity: surface area, length, basal width, orientation
 - ➔ Myocardial surface
 - ➔ Epicardial wall center

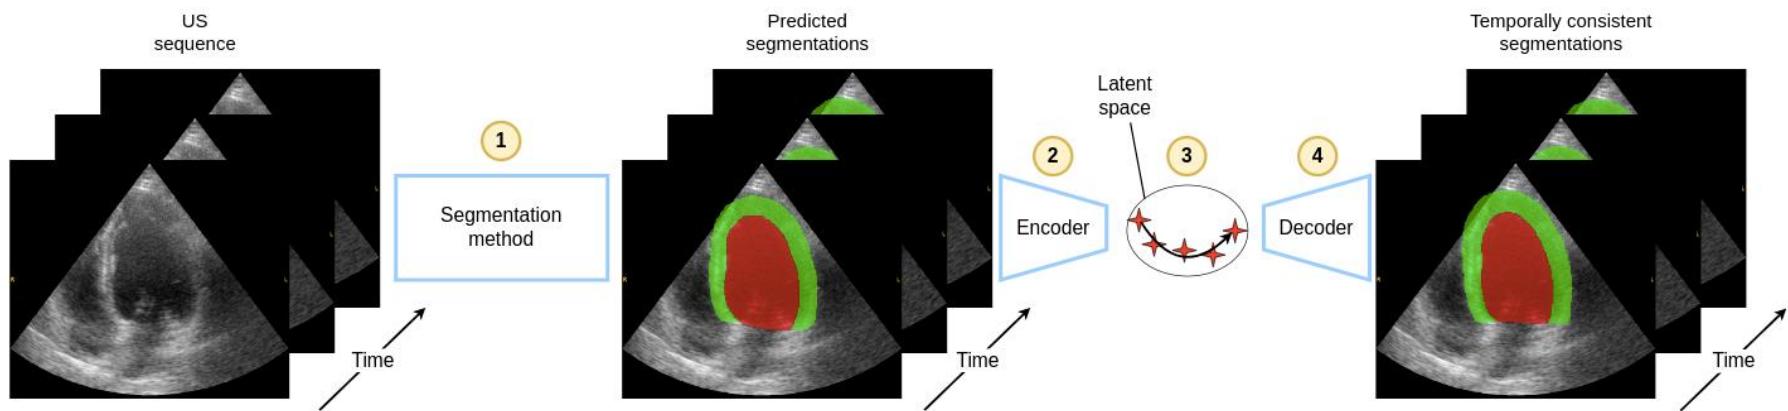


Shape representation




Cardiac segmentation with temporal consistency

- ▶ Post-processing to ensure temporal consistency



[Painchaud, IEEE TMI, 2022]


Cardiac segmentation with temporal consistency




Cardiac segmentation with temporal consistency

- ▶ Some post-processing examples

Original segmentation




Post-processed segmentation




Cardiac segmentation with temporal consistency

- ▶ Some post-processing examples

Original segmentation



Post-processed segmentation



Hands-on session

Hands-on session

- ▶ <https://olivier-bernard-creatis.github.io//teaching/>

Olivier Bernard Research Publications Talks Teaching Blog Posts CV



Teaching

Diffusion model

Course, INSA, university of Lyon, telecommunications department, 2024

Course on diffusion model. The following aspects are covered: markov chain, denoising diffusion probabilistic model, conditioning, latent diffusion model.

[[pdf-fr](#)] [[pdf-en](#)] [[code 1](#)] [[code 2](#)] [[post](#)]

Transformers

Course, INSA, university of Lyon, electrical department, 2024

French course on transformers. The following aspects are covered: tokenisation, positional embedding, encoding blocs, self-attention module, multi-head attention, vision transformer.

[[pdf-fr](#)] [[pdf-en](#)] [[code 1](#)] [[post](#)]

Variational Auto-encoders

Course, INSA, university of Lyon, electrical department, 2024

Course on variational auto-encoder networks. The following aspects are covered: fundamental concepts, Kullback-Liebler divergence, variational inference, VAE architectures.

[[pdf-fr](#)] [[pdf-en](#)] [[code 1](#)] [[code 2](#)] [[code 3](#)] [[code 4](#)] [[post](#)]

Diffusion models

Diffusion models


- ▶ Best current methods for synthetic image generation
- ▶ Allows generating images in a *conditioned* form
- ▶ Many software solutions, such as Midjourney, DALL-E

An Asian girl in ancient coarse linen clothes rides a giant panda and carries a wooden cage. A chubby little girl with two buns walks on the snow. High-precision clothing texture, real tactile skin, foggy white tone, low saturation, retro film texture, tranquil atmosphere, minimalism, long-range view, telephoto lens



What is the purpose of diffusion models?

► Family of diffusion networks




Denoising Diffusion
Probabilistic models



Intuition behind diffusion models

Variational Auto-encoders

- Completeness is expressed as a **soft constraint** !
 - $\mathcal{N}(g(x), \text{diag}(h(x)))$ and $\mathcal{N}(0, I)$ should remain close in terms of distributional distance




Sampling from the latent space $\mathcal{N}(0, I)$ does not guarantee to obtain a reconstructed image from the target distribution

Variational Auto-encoders


► Illustration from Mednist dataset

- (train,valid,test) = (1491,373,223)
- Input image size: 48x48 / latent space $K=432$ (compression factor around 5)



Variational Auto-encoders


- ▶ Linear interpolation between two real images



Variational Auto-encoders

- ▶ Sampling directly from the latent space

$$z \in \mathbb{R}^{(K)} \text{ with } z_i \sim \mathcal{N}(0, I)$$



A soft constraint on the latent space to remain close to $\mathcal{N}(0, I)$ is not sufficient to build generative models that effectively learn a target distribution

The denoising diffusion probabilistic models

DDPM


All the mathematics are described in the following blog

<https://creatis-myriad.github.io/tutorials/2023-11-30-tutorial-ddpm.html>


Basic idea of denoising diffusion model

How can a hard constraint be enforced to ensure a direct transformation from the latent space (modeled as a Gaussian) to the target distribution?

► Noising process




► Learning of the denoising process



Noising process (forward diffusion process)

- Modeled as a sequence of normal distributions (Markov chain process)

$$q(x_t | x_{t-1}) = \mathcal{N} \left((\sqrt{1 - \beta_t}) x_{t-1}, \beta_t \mathbf{I} \right)$$




$$q(x_t | x_{t-1}) = (\sqrt{1 - \beta_t}) x_{t-1} + \beta_t \mathbf{I} \cdot \epsilon_t$$

Noising process (forward diffusion process)


- Modeled as a sequence of normal distributions (Markov chain process)

$$q(x_t | x_{t-1}) = \mathcal{N} \left((\sqrt{1 - \beta_t}) x_{t-1}, \beta_t \mathbf{I} \right)$$




- β_t : variance varying over the iterative process from 0 to 1

if $\beta_t = 0$, then $q(x_t | x_{t-1}) = x_{t-1}$
if $\beta_t = 1$, then $q(x_t | x_{t-1}) = \mathcal{N}(0, \mathbf{I})$




Noising / denoising processes

► Noising process




$$\begin{cases} \alpha_t = 1 - \beta_t \\ \bar{\alpha}_t = \prod_{k=1}^t \alpha_k \\ \epsilon_t \sim \mathcal{N}(0, \mathbf{I}) \end{cases}$$

► Denoising process



Training procedure


- ▶ Choose a random step $t \in \{1, \dots, T\}$
- ▶ Train a U-Net model to predict the noise pattern ϵ_θ to remove from x_t



Architecture


Standard U-Net with attention layers and position encoding to integrate temporal information

→ Integration of t is necessary because the added noise varies over time




Architecture

→ Attention layer



Inference: generation of synthetic data



- ▶ Generate a random image $x_T \sim \mathcal{N}(0, I) \in \mathbb{R}^{N \times M}$
- ▶ At each step from T to 0, use the U-Net model to compute

$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) + \sigma_t \epsilon$$

U-Net

with $\epsilon \sim \mathcal{N}(0, \mathbf{I})$

Mathematical formalism

Mathematical formulation

► Useful notations

$$q(x_1, \dots, x_T | x_0) = q(x_1 | x_0) q(x_2 | x_1, x_0) \cdots q(x_T | x_{T-1}, \dots, x_0)$$



Complete forward
process

$$q(x_1, \dots, x_T | x_0) = q(x_1 | x_0) q(x_2 | x_1) \cdots q(x_T | x_{T-1})$$



Markov chain

$$q(x_{1:T} | x_0) = q(x_1 | x_0) q(x_2 | x_1) \cdots q(x_T | x_{T-1})$$



Compact
reformulation

$$q(x_{1:T} | x_0) = \prod_{t=1}^T q(x_t | x_{t-1})$$

$$q(x_{1:T} | x_0) = \prod_{t=1}^T q(x_t | x_{t-1})$$

Complete forward process

$$p_\theta(x_{0:T}) = p(x_T) \prod_{t=1}^T p(x_{t-1} | x_t)$$

Complete reverse process

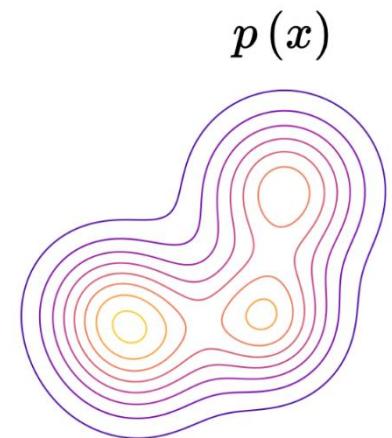
Mathematical formulation

► Optimization process

→ Maximization of $\log(p_\theta(x))$ / Minimization of $-\log(p_\theta(x))$

$$-\log(p_\theta(x))$$

$$-\log(p_\theta(x)) = -\log \left(\int p_\theta(x_{0:T}) dx_{1:T} \right)$$



← Marginal distribution

$$-\log(p_\theta(x)) = -\log \left(\int \frac{q(x_{1:T} | x_0)}{q(x_{1:T} | x_0)} p_\theta(x_{0:T}) dx_{1:T} \right)$$

$$-\log(p_\theta(x)) = -\log \left(\mathbb{E}_{q(x_{1:T}|x_0)} \left[\frac{p_\theta(x_{0:T})}{q(x_{1:T} | x_0)} \right] \right)$$



Expectation
reformulation




Jensen's inequality

$$-\log(p_\theta(x)) \leq -\mathbb{E}_{q(x_{1:T}|x_0)} \left[\log \left(\frac{p_\theta(x_{0:T})}{q(x_{1:T} | x_0)} \right) \right]$$

Mathematical formulation

► Evidence lower bound (ELBO)

$$-\log(p_\theta(x)) \leq -\mathbb{E}_{q(x_{1:T}|x_0)} \left[\log \left(\frac{p_\theta(x_{0:T})}{q(x_{1:T} | x_0)} \right) \right] \quad \text{ELBO}$$



→ Minimization of the ELBO

$$-\mathbb{E}_{q(x_{1:T}|x_0)} \left[\log \left(\frac{p_\theta(x_{0:T})}{q(x_{1:T} | x_0)} \right) \right]$$

⋮

$$-\mathbb{E}_q \left[D_{KL} (q(x_T | x_0) \| p_\theta(x_T)) + \sum_{t>1} D_{KL} (q(x_{t-1} | x_t, x_0) \| p_\theta(x_{t-1} | x_t)) - \log (p_\theta(x_0 | x_1)) \right]$$



No parameter to
be learned




Very small

Mathematical formulation

► ELBO minimization

$$\mathcal{L} = -\mathbb{E}_q \left[\sum_{t>1} D_{KL} (q(x_{t-1} | x_t, x_0) \| p_\theta(x_{t-1} | x_t)) \right]$$



→ Exploitation of the Gaussian properties of the forward process and modeling of the reverse process using Gaussian distribution

$$q(x_{t-1} | x_t, x_0) = \mathcal{N} \left(\tilde{\mu}_t, \tilde{\beta}_t \right) \quad p_\theta(x_{t-1} | x_t) = \mathcal{N} (\mu_\theta, \sigma_\theta)$$

→ Reformulation

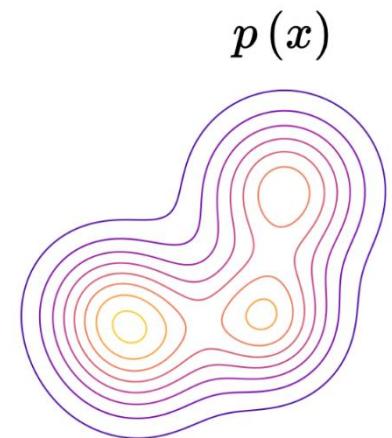
$$D_{KL} (q(x_{t-1} | x_t, x_0) \| p_\theta(x_{t-1} | x_t)) = -\frac{1}{2\sigma_t^2} \|\tilde{\mu}_t - \mu_\theta\|^2$$

$$\mathcal{L} = \mathbb{E}_q \left[\sum_{t>1} \frac{1}{2\sigma_t^2} \|\tilde{\mu}_t - \mu_\theta(x_t, t)\|^2 \right]$$

Mathematical formulation

► ELBO minimization

$$\mathcal{L} = \mathbb{E}_q \left[\sum_{t>1} \frac{1}{2\sigma_t^2} \|\tilde{\mu}_t - \mu_\theta(x_t, t)\|^2 \right]$$



➔ Expressions of means

$$\tilde{\mu}_t = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_t \right) \quad \mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right)$$

$$\mathcal{L} = \mathbb{E}_q \left[\sum_{t>1} \frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1-\bar{\alpha}_t)} \|\epsilon_t - \epsilon_\theta(x_t, t)\|^2 \right]$$

➔ Simplifications

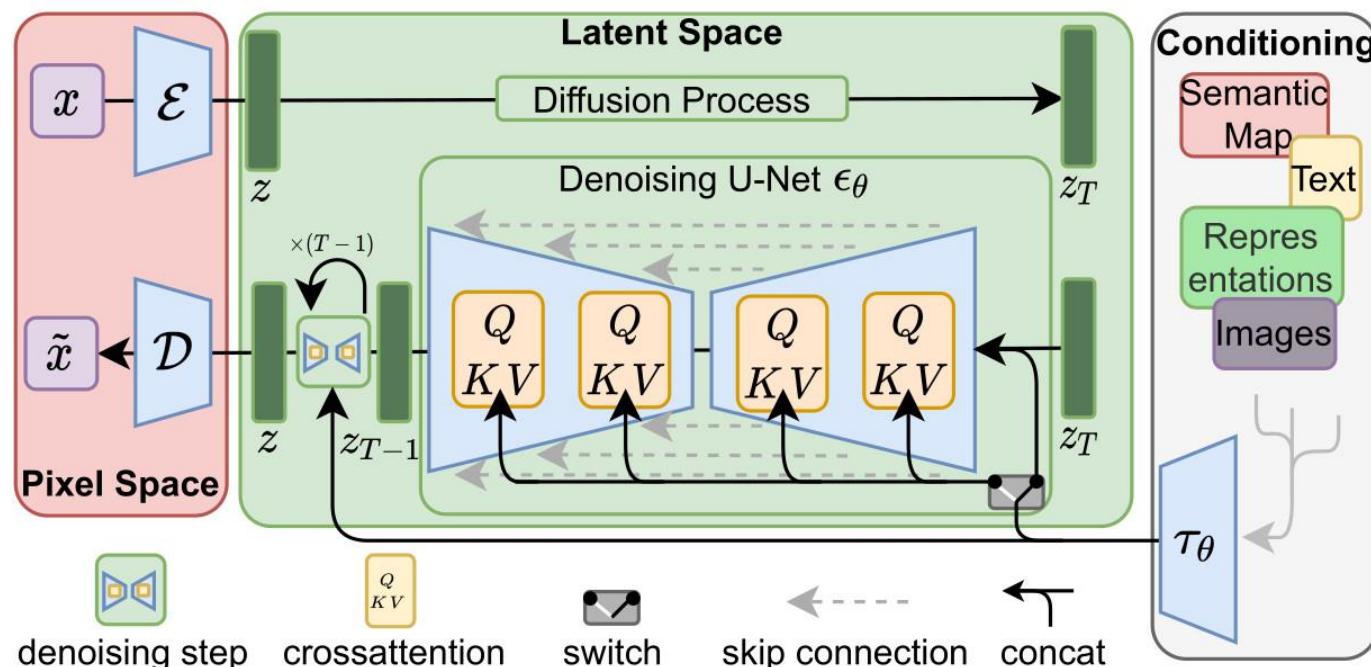
$$\mathcal{L} = \mathbb{E}_{q,t} [\|\epsilon_t - \epsilon_\theta(x_t, t)\|^2]$$

Practical application

Latent diffusion models

Latent diffusion model (LDM)

- ▶ VAE is learned independently of DDPM and its architecture is fixed
 - ▶ Efficiently reduce the dimensionality of the input space
 - ▶ Efficiently initiate the Gaussian diffusion process
- ▶ LDM architecture



Latent diffusion model (LDM)

► Properties

Parameters	LDM – 256 × 256
z dimensions	64 × 64 × 3
Diffusion steps	1000
Noise scheduler (β_t)	linear
Number of parameters	274 Million
Channels	224
Channel multiplier	1, 2, 3, 4
Levels for attention	2, 3, 4
Number of head	1
Batch size	48
Iterations	410 k
Learning rate	9.6 e⁻⁵

Latent diffusion model (LDM)


- ▶ Random generation of synthetic images *without conditioning* learned from the CelebA-HQ database

Random samples on the CelebA-HQ dataset






Latent diffusion model (LDM)

- ▶ Random generation of synthetic images *with conditioning on the class* learned from the ImageNet database



Latent diffusion model (LDM)

- ▶ Random generation of synthetic images *with conditioning on masks* learned from the Flickr-landscapes database




Latent diffusion model (LDM)


- ▶ Random generation of synthetic images *with conditioning on text* learned from LAION-400M database
 - ➔ Using the BERT tokenizer
 - ➔ This model has over 1.45 billion parameters!

'A painting of the last supper by Picasso.'




Medical applications





Diffusion models for data augmentation

- ▶ Synthetic dataset generation for brain MR volumes [Walter et al., MICCAI workshop 2022]
- ▶ UK Biobank dataset
 - ▶ 3D MR volumes (T1w)
 - ▶ Training: 31,740 patients
 - ▶ with covariates: age (44 to 82 years), gender (53% women), brain structure volumes
 - ▶ Quality of synthetic data measured using FID: Fréchet Inception Distribution



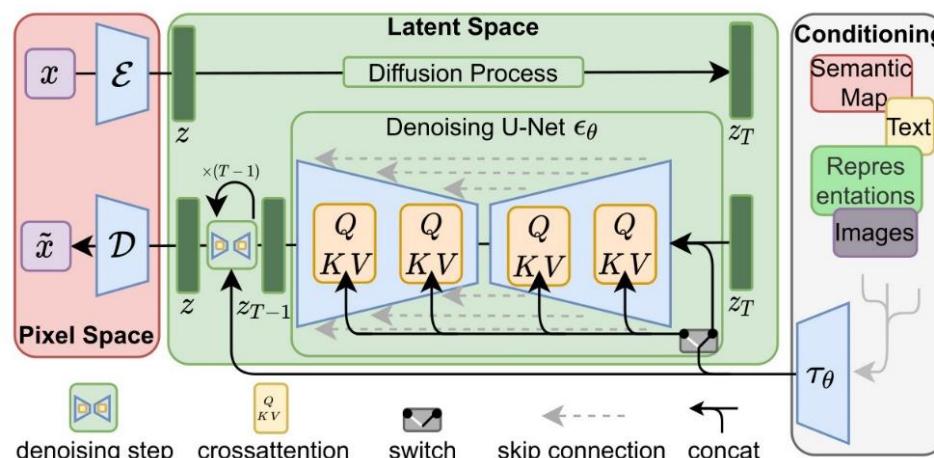
Diffusion models for data augmentation

► VAE

- ▶ 3D convolutions
- ▶ Latent space dimension: $20 \times 28 \times 20$

► DDPM

- ▶ 3D convolutions
- ▶ T=1000 time steps
- ▶ Conditioning: vector encoding each covariable




Diffusion models for data augmentation

► Results

- FID: generated from 1,000 samples drawn from each of the two distributions to be compared



	FID ↓
LSGAN	0.0231
VAE-GAN	0.1576
LDM	0.0076
Real images	0.0005




Diffusion models for data augmentation

► Results

- ▶ SynthSeg model was used to automatically measure brain volumes from synthetic data
- ▶ A 3D CNN trained from the UK biobank was used to automatically predict the age from the synthetic data




- ▶ Synthetic dataset of 100,000 human brain was generated and made publicly available with the conditioning information
- ▶ Promote data sharing with privacy guarantees



Diffusion models for data augmentation

- ▶ Synthetic dataset generation for brain MR volumes [El-Allaly et al., Eusipco 2025]
- ▶ Set of public datasets (IBSR, OASIS, IXI, Kirby)
 - ▶ 3D MR volumes (T1w, T2w, FLAIR)
 - ▶ Training: 40,000 axial MRI slices
 - ▶ with **textual description** generated from atlas registration + [...]




Diffusion models for data augmentation

- ▶ Synthetic dataset generation for brain MR volumes [El-Allaly et al., Eusipco 2025]
- ▶ Set of public datasets (IBSR, OASIS, IXI, Kirby)
 - ▶ 3D MR volumes (T1w, T2w, FLAIR)
 - ▶ Training: 40,000 axial MRI slices
 - ▶ with **textual description** generated from atlas registration + metadata + natural language description using template-based text synthesis

ID Labels
2 white matter of left cerebral hemisphere
3 cortex of left cerebral hemisphere
4 left lateral ventricle
5 temporal horn of left lateral ventricle
6 ...

Volume information expl. IXI to IBSR
modality T2
mov_gender F
mov_age 71
fix_gender M
fix_age 37
Sx mm 1.25
Sy mm 1
Sz mm 1

Labeled image




Diffusion models for data augmentation

- ▶ Text encoder
 - ▶ WordPiece tokenizer
 - ▶ BiomedBERT pre-trained model
 - ▶ Number of tokens processed : 512
 - ▶ Token size (text embedding dimension): 768
- ▶ VAE
 - ▶ Pre-trained from 40,000 axial MRI slices
 - ▶ 2D convolutions
 - ▶ Input image dimension: 256 x 256 x 1
 - ▶ Latent space dimension: 64 x 64 x 1
 - ▶ Training time (GPU 32 GB): 40 hours

Diffusion models for data augmentation

► DDPM

- ▶ 2D convolutions
- ▶ T=1000 time steps
- ▶ Conditioning: set of token encoding the text description of size 512 x 768
- ▶ Input size: 256 x 256 x 1 / Latent space 64 x 64 x 1 / training time (32 GB): 14 days



Diffusion models for data augmentation


► Results

- FID: generated from 1,000 samples drawn from each of the two distributions to be compared


	FID ↓
Train/val	16.9
Test	17.9

"MRI image of the head of a woman of 21 years old in **T1-weighting**, with axial section. We observe slice n°140 beginning from the base of the skull. The observed details include the white matter of left cerebral hemisphere of volume 2728.75 mm³, the cortex of left cerebral hemisphere of volume 4366.25 mm³, the left lateral ventricle of volume 267.50 mm³, the left thalamus of volume 202.50 mm³, the left caudate nucleus of volume 162.50 mm³, the left putamen of volume 180.00 mm³, the cerebrospinal fluid of volume 7.50 mm³, the white matter of right cerebral hemisphere of volume 3017.50 mm³, the cortex of right cerebral hemisphere of volume 4041.25 mm³, the right lateral ventricle of volume 250.00 mm³, the right thalamus proper of volume 217.50 mm³, the right caudate nucleus of volume 165.00 mm³, the right putamen of volume 135.00 mm³ and the right thin cerebral white matter of volume 685.00 mm³."

Input prompt



Synthetic image with key word T1-weighted



Synthetic image with key word T2-weighted

Diffusion models for data augmentation


► Results

- FID: generated from 1,000 samples drawn from each of the two distributions to be compared


	FID ↓
Train/val	16.9
Test	17.9

"MRI snapshot with **T1-weighting** describing the skull of a woman of age of 52 years old, of axial section n°110 from the base of the skull. We distinguish clearly: the white matter of left cerebral hemisphere of volume 171.25mm^3 , the cortex of left cerebral hemisphere of volume 1291.25mm^3 , the white matter of left hemisphere of cerebellum of volume 550.00 mm^3 , the left cerebellar cortex of volume 1483.75 mm^3 , the fourth ventricle of volume 97.50 mm^3 , the brainstem of volume 582.50 mm^3 , the white matter of right cerebral hemisphere of volume 222.50 mm^3 , the cortex of right cerebral hemisphere of volume 1181.25 mm^3 , the white matter of right hemisphere of cerebellum of volume 593.75 mm^3 , the right cerebellar cortex of volume 1413.75 mm^3 and the right thin cerebral white matter of volume 302.50 mm^3 ."


Input prompt



Synthetic image with
key word T1-weighted




Synthetic image with
key word T2-weighted



Diffusion models for image segmentation

- ▶ Segmentation of tumors from MR images [Wolleb et al., MIDL 2022]
- ▶ BRATS2020 dataset
 - ▶ 4 different MR sequences per patient (T1, T2, T1ce, FLAIR)
 - ▶ Training: 332 patients with 3D volumes sequences => 16,998 2D images
 - ▶ Testing: 37 patients with 3D volumes sequences => 1,082 2D images





4 MR inputs per patient (T1, T2, T1ec, FLAIR)



Mask output


Diffusion models for image segmentation

- ▶ Learn the underlying distribution of tumor segmentation masks



Diffusion models for image segmentation

- ▶ Conditioning with the 4 MR images using concatenation scheme



Diffusion models for image segmentation

- At inference time: modelling of the segmentation uncertainty




Illustration taken from <https://www.youtube.com/watch?v=US9CzPrT2H8>

Diffusion models for image segmentation

► Results






Illustration taken from <https://arxiv.org/pdf/2112.03145>



Diffusion models for anomaly detection

- ▶ Anomaly detection from MR images [Wolleb et al., MICCAI 2024]
- ▶ BRATS2020 dataset
 - ▶ 4 different MR sequences per patient (T1, T2, T1ce, FLAIR)
 - ▶ Training: 332 patients with 3D volumes sequences => 16,998 2D images
 - ▶ 5,598 healthy 2D slices (without tumor) / 10,607 disease 2D slices



4 MR inputs per patient (T1, T2, T1ec, FLAIR)





Mask output

Diffusion models for anomaly detection

► General idea

Unpaired image-to-image translation




How to preserve spatial anatomical information using a diffusion process?

Diffusion models for anomaly detection

► Denoising Diffusion Implicit Models (DDIM)


- Reformulation of the diffusion process
- Remove the random component $\sigma_t \epsilon$
- Make the diffusion process deterministic



Diffusion models for anomaly detection

► Main algorithm – part 1

- Train a classical DDPM on the dataset containing healthy and disease images
- Train a classifier network C to predict the class label (healthy vs disease) from any noisy images x_t



Diffusion models for anomaly detection

► Main algorithm – part 2

- Use DDIM process
- Compute the gradient of the classifier to guide the removing of anomaly regions





Illustration taken from <https://arxiv.org/pdf/2203.04306>


Diffusion models for anomaly detection

- ▶ Result on an image with a tumor



Diffusion models for anomaly detection

- ▶ Result on an image without any tumor



That's all folks

Variation Auto Encoder framework

► Key idea

- ➔ Generating sample z according to $p(z|x)$, implies that z has been generated by images from the original data distribution $p(x)$
- ➔ If we can reconstruct vectors back into images, we will effectively generate new samples from our original data distribution
- ➔ We need to know the latent distribution, which is assumed to be a normal distribution
- ➔ This allows to compute the likelihood $p(x|z)$
- ➔ The only unknown remains the true posterior $p(z|x)$
- ➔ Thanks to variational inference, we approximate it using a Gaussian distribution $q(z|x)$
- ➔ This Gaussian will have parameters μ and σ that we need to learn, which is an optimization process known as variational Bayes
- ➔ We will train an encoder to estimate these parameters μ and σ from the images
- ➔ Then we used a decoder to reconstruct images from the latent variables that are sampled from the approximate posterior